

Operating systems

ELTE IK.

Dr. Illés Zoltán

zoltan.illes@elte.hu

We have talked earlier...

- ▶ **Progress of Operating systems**
 - Computer – Op.system generations
- ▶ **Notions of Operating system, structures**
 - Client–server model, system call
- ▶ **Files,directories, filesystems**
 - Hardware architecture
 - Logical architecture
 - FAT, UNIX, NTFS,...
- ▶ **Processes, process communications**
 - Race conditions, critical sectionmodel

What comes today...

- ▶ **Process communication**
 - Monitors
 - Message sending
- ▶ **Classical IPC problems**
 - Case of eating philosophers
- ▶ **Scheduling of processes**
 - Theories, implementations
 - Thread scheduling

Is there any problem with semaphores?

- ▶ No...but small mistyping might cause big problems.
 - E.g. Let's change 2 rows (red), when bakery is full, the baker is stopped by semaphor empty, at the same time the customer is stopped by free, therefore both processes are stopped! (deadlock)

```
void baker() /*original recipe
*/
{
    int bread;
    while (1)
    {
        bread=baker_bake();
        down(&empty);

        down(&free);
        to_shelf(bread);
        up(&free);
        up(&full);
    }
}
```

```
void baker() /*modified recipe*/
{
    int bread;
    while (1)
    {
        bread=baker_bake();
        down(&free); /*changed part*/
        down(&empty);
        to_shelf(bread);
        up(&free);
        up(&full);
    }
}
```

Do we have any problem with semaphors?

- ▶ Generally no! But as we saw, a little mistake causes program error! (Big error!).
- ▶ Also the changing two up and down instructions can cause same problem!
- ▶ Other reason: if we forget it...
- ▶ Do we have any better thing to do?
 - At Kernel level no.

Monitors

- ▶ Brinch Hansen (1973), Charles Anthony Richard Hoare (1974) suggested a higher level construction supported by languages.
- ▶ This construction is called: monitor
 - It is similar to a class definitions

```
Monitor critical_zone
    Integer shelf[];
    Condition c;
    Procedure baker(x);
    ...
End;
Procedure customer(x);
...
End;
End monitor;
```


Monitor properties

- ▶ We can define inside a monitor procedures, data structures.
- ▶ At the same time only 1 process can enter into a monitor!
- ▶ This is supported by programming language compiler.
 - If a process calls a method inside a monitor, first it checks whether there is any other active process inside the monitor?
 - If yes, the process will be held, suspended.
 - If not, the process enters into the monitor.

Monitor implementation

- ▶ Using Mutex
- ▶ A customer has no concrete idea about implementations, but it is not needed.
- ▶ Result: a more safety mutual exclusion
- ▶ Only we have one problem: what about if a process can't continue his work inside a monitor?
 - E.g: a baker can't bake bread, because the bakery shelves are full!
- ▶ Solution: lets create a condition variable
 - We can do 2 operation on condition variable: wait, signal

A Solution of the producer–consumer problem with monitors I.

N element

```
monitor Baker–Customer
  condition full, empty;
  int number;
  bread_to_shelf(bread element)
  {
    if (number==N) wait(full);
    shelf(element);
    number++;
    if (number==1) signal(empty);
  }
  bread bread_from_shelf()
  {
    if (number==0) wait(empty);
    bread element=bread_from_shelf();
    number--;
    if (number==N-1) signal(full);
    return element;
  }
end monitor
```

The process of Pék&Vásárló

```
baker()
{
    while(1)
    {
        bread new;
        new=bread_bake();
        Baker-Customer.bread_to_shelf (new);
    }
}
customer()
{
    while(1)
    {
        bread new_bread;
        new_bread=Baker-Customer.bread_from_shelf();
        eat(new_bread);
    }
}
```

Other solutions

- ▶ The previous example based on a theoretical language: Pidgin Pascal
- ▶ In C there is no monitor
 - In C++ there is monitor, also wait, notify
- ▶ Java:
 - Synchronized methods
 - No condition variable, but there are wait & notify
- ▶ C#
 - Monitor class
 - Enter, TryEnter, Exit, Wait, Pulse (this same as notify)
 - Lock keyword
 - Sample: VS2008 solution, monitor project.

Do we have any problem with monitors?

- ▶ Hm,...no, we have not!
- ▶ It is more safety than a simple semaphor.
- ▶ It works with one or more CPU, but only with a common memory space!
- ▶ If a CPU has a standalone memory this solution does not fit, does not work!

Message sending

- ▶ Processes typically uses two general method:
 - Send(to, message)
 - Receive(from, message)
 - Source may be optional!
- ▶ These are system calls and not language constructions!
- ▶ Receipt message: If a sender wants to know about accepting message, the receiver process sends back a receipt!
 - If a sender does not get the receipt, it will send again the message!
 - If the receipt will be lost, the sender also will send again!
 - A message contain a number which holds the message number too to distinguish(mark) messages .

A Solution of the producer–consumer problem with message sending. I.

► The producer (baker) process:

```
#define N 100          // size of shelves
void baker()          // baker process
{
    int bread;         // „bread” variable, storage
    message m;         // message storage place
    while(1) // baker works all time
    {
        bread= bake_bread();
        receive(customer,m);    // we wait a message from
                                // the customer
                                // empty message in m
        m=create_message(bread);
        send(customer,m); // send the bread to the customer
    }
}
```


A Solution of the producer–consumer problem with message sending. II.

► Consumer process :

```
void customer()          // consumer (customer) method
{
    int bread;           // „bread” variable
    message m;           // the place for the message
    int i;
    for(i=0;i<N;i++) send(baker,m);
                        //setup N empty space for bred
                        //sending to baker
    while(1) // the shopping is going on all the time too
    {
        receive(baker,m);      // buying a bread
        bread=message_outpack(m);
        send(baker,m); // sending back empty store
        eat_bread (bread);
    }
}
```

Summary of message sending

- ▶ We have to create temp. place for messages on both places. (like mail store)
- ▶ We can omit this, in this case the receive statement will be blocked before send command!
 - Trysting strategy.
 - Minix 3 also uses this one!
 - Messages are the same as pipes, but inside pipes there is no any delimiters!
- ▶ The message sending is a general technique in parallel systems.

Classic IPC problems I.

- ▶ Eating philosophers case:
 - They need 2 forks to eat spaghetti!
 - All philosophers can access only forks near his own plate.
 - They are thinking-eating...
 - Lets create a program to solve this problem!



Solution I.

- ▶ It has a problem, all processes can be in deadlock if all gets the lefts fork and will wait for the right one!.
- ▶ If somebody puts back left fork and again try to get both also not good, because all processes do this! (Starvation)

```
Void philosopher (int i)
{
    while(1)
    {
        thinking();
        need_fork(i); // left fork
        need_fork ((i+1)%N); //right
        eating();
        dont_need_fork (i);
        dont_need_fork ((i+1)%N);
    }
}
```

Solution II.

```
int s[5];           // values: eat, hungry, thinking
semaphore safe_s=1; // a sign for using s array
semaphore philo[5]={0,0,0,0,0}; // 1 semaphore to each philosopher, 0=not
permitted
Void philosopher(int i)
{
    while(1) {
        thinking();
        down(safe_s); // only i modify s[]
        s[i]=hungry; //
        if (s[left]!=eat && s[right]!=eat) // whether the two neighbour's fork are free?
            { s[i]=eat; up(philo[i]); }; // i eat, philo[i] shows free
        up(safe_s); // other may access s[]
        down(philo[i]);
        // it is blocked, if there is no 2 forks, if the i. Philosopher is not eating
        spaghetti_eating();
        down(safe_s); // finished eating, again protect s because i modify it
        s[i]=thinking;
        if (s[left]==hungry && s[left--]!=eat) { s[left]=eat; up(philo[left]); }
        if (s[right]==hungry && s[right++]!=eat) { s[right]=eat; up(philo[right]); }
        up(safe_s);
    }
}
```


Solution III.

- ▶ Lets 5 forks semaphore[] for each fork.
- ▶ Max semaphore
- ▶ Sample to access restricted resources.

```
Int N=5;
semaphore forks[]={1,1,1,1,1}; //all free
semaphore max=4; //max 4 fork used in
the same time
void philozophus(int i)
{
    while(1)
    {
        thinking();
        down(max);
        down(forks[i]); // left fork
        down(forks[(i+1)%N]); //right
        eating();
        up(forks[i]);
        up(forks[(i+1)%N]);
        up(max);
    }
}
```


Readers–Writer problem

- ▶ The database can read by more than 1 process, but only 1 can write:

```
// writing process
semaphore database=1;
semaphore mutex=1;
int rc=0;
Void writer()
{
    while(1)
    {
        do_something();
        down(database); // critical
        write_to_database();
        up(database);
    }
}
```

```
Void reader()
{
    while(1)
    {
        down(mutex);
        rc++;
        if (rc==1) down(database);
        up(mutex);
        Read_from_database();
        down(mutex); // critical
        rc--;
        if (rc==0) up(database);
        up(mutex);
        working_on_data();
    }
}
```

Scheduling

- ▶ We have seen before that several processes may run „parallel” to each other.
- ▶ Only 1 at a time can run.
- ▶ Which one should run?
- ▶ The scheduler makes the decision
- ▶ The decision is based on the scheduling algorithm

The I/O need of processes

- ▶ Typically a process may have two different types of work:
 - It is counting...
 - I/O need, wants to read or write data from/to the periphery
- ▶ CPU-bound process
 - It is working (count) a lot, it does not wait to much for I/O
- ▶ I/O-bound process
 - Works only for a small amount of time it is waiting for I/O long

When do we have to switch to a new process?

- ▶ There is a process (context) switch:
 - If a process ended
 - If a process' state becomes blocked (because I/O or semaphore)
- ▶ Usually there is a context switch:
 - A new process is created
 - I/O interrupt occurs
 - After an I/O interrupt, a blocked process, the process which waited for it typically may continue its running.
 - Timer interrupt
 - Preemptive scheduling
 - Non-preemptive scheduling

Groups of scheduling

- ▶ Representatives for each system:
 - Fairness – everybody can use CPU
 - The same rules are valid for everybody
 - Balance –Everybody should get the same loading
- ▶ Batched systems
 - throughput, turnaround time, CPU utilization
- ▶ Interactive systems
 - Response time, proportionality to the user's expectations
- ▶ Real-time systems
 - To keep deadlines, to avoid data loss, quality corruption

Scheduling in batch systems I.

- ▶ Ranking scheduling, it is non-preemptive
 - First Come First Served – (FCFS)
 - The process runs till it stops or becomes blocked.
 - If it blocks, it goes to the end of the queue.
 - The processes are in a fair, simple, chained list.
 - Disadvantage: I/O-bound processes are very slow.
- ▶ The shortest job runs first, non-preemptive (SJB)
 - We have to know the running time ahead
 - It is optimal, if everybody is known at the beginning

Scheduling in batch systems II.

- ▶ The process with the shortest remaining time runs
 - Preemptive, monitoring at each entering.
- ▶ Three level scheduling
 - Inlet scheduler
 - It lets the task into the memory in rotation.
 - Disk scheduler
 - If it lets in to much processes and memory becomes full then they all have to be written to the disk and back.
 - It runs rarely.
 - CPU scheduler
 - We may choose from the previously mentioned algorithms.

Scheduling in interactive systems I.

► Round Robin

- Everybody gets a time quantum, at the end of it or in the case of blocking comes the following process
- At the end of the quantum the next process in the round list will be the actual
- Fair, simple
- We may store the processes (features) in a list and we go round and round all the time.
- There is only one question: How big should be the quantum?
 - The process switch needs some amount of time.
 - Small quantum → a lot of CPU time for switching
 - Too big quantum → the interactive user feels it too slow (keyboard handling)

Scheduling in interactive systems II.

► Priority scheduling

- Importance, priority is shown
 - Unix: 0–49 → not preemptive (kernel) priority
 - 50–127 → user priority
- The process with the highest priority may run
 - Modifying: dynamical priority to avoid starvation
- Usage of priority classes
 - Usage of Round Robin among the same class processes
 - Must modify the priority of the processes, because low priority processes get CPU rarely.
 - Typically at each 100. quantum the priorities are recounted
 - Typically the high priority processes become to lower level, then comes RR. At last the original classes will be built up again.

Scheduling in interactive systems III.

► Multiple queues

- Also have priorities and uses RR
- At the highest level each quantum gets 1 time quantum.
- The next 2, then 4, 8, 16, 32, 64.
- If the time of the highest process is used up it goes down with one level.

► The shortest process next

- Though we do not know the remaining execution time we have to estimate it!
- Aging, weighted average for the time quantum.
 - $T_0, T_0/2 + T_1/2, T_0/4 + T_1/4 + T_2/2, T_0/8 + T_1/8 + T_2/4 + T_3/2$

Scheduling in interactive systems IV.

- ▶ **Guaranteed Scheduling**
 - Each active process gets proportional CPU time.
 - Must be registered how many time was used up for a process and if somebody has got less time it „goes” forward.
- ▶ **Lottery scheduling**
 - Similar to the previous one, except the processes gets some lottery ticket. That process can run which has got the pulled out one.
 - It is easy to support proportional CPU time, useful e.g. at video servers
- ▶ **Fair-share scheduling**
 - Let's pay attention to the user as well! Similar to guaranteed one – only it refers to the users

Scheduling in real-time systems I.

- ▶ What is a real-time system?
 - The time is an important actor. We must guarantee to give a response within the deadline.
 - Hard Real Time, absolute, not modifyable deadlines.
 - Soft Real Time (tolerant), there are deadlines, but a small difference is tolerable.
 - The programs are divided into several smaller process.
 - In the case of an outer event, have to give a response within the deadline.
 - Schedulable: if the response CPU time sum of n event is ≤ 1 .
- ▶ Unix, Windows are real-time systems?

Scheduling ideas, implementation

- ▶ Frequently there are child processes in the system.
- ▶ It is not sure that the priority of the parent and the child process must be the same.
- ▶ Typically the kernel uses priority scheduling (+RR)
 - It warrants a system call with which the parent may give the priority of the child
 - Kernel schedules – the user process modify the priority (nice)

Threadscheduling

► User level threads

- The Kernel does not know about it, the process gets a quantum, within the thread scheduler makes the decision who should run
- Quick switch between the threads
- It is possible to use application dependent thread scheduling

► Kernel level threads

- Kernel knows the threads, Kernel decides which process which thread should run.
- Slow switch, between the switch of two threads needs a full context switch
- This is also noticed.

Thanks for your
attention!

zoltan.illes@elte.hu